

IDS PROJECT REPORT

**Implementing ML Classification Algorithm for
Predicting Whether the Income exceeds a given Income based on
census data**

<https://archive.ics.uci.edu/ml/datasets/adult>

Team Members:

Swati Mor	: 21ucc101
Arpit Gupta	: 21ucs030
Arpit Jain	: 21ucs031
Elishben Manojbhai Baraiya	: 21ucs077

Course Instructors:

Dr. Alope Datta
Dr. Subrat Das
Dr. Lal Upendra Pratap Singh

Problem Statement

We are given an Excel dataset that has 15 columns and 32561 rows.

We are trying to predict whether the income of an adult will exceed 50k per year or not by developing a supervised machine learning model.

To achieve it we need to follow the data science/ML life cycle. Starting with data collection/extraction(which we have already collected). After that we need to follow following steps:

1. Data pre-processing and its visualization and explain all the inferences we got from our data.
2. Decide what ML Classification Algorithms to use and why
3. Implementing those algorithms
4. Output the result of the testing set and its visualization

Python has many libraries which makes the above steps more easy and efficient. Below is the list of some of libraries we would be using:

- `scikit_learn`
- `matplotlib`
- `seaborn`
- `numpy`
- `pandas`

Implementation

1.Importing the libraries and loading the dataset

(i) Importing Libraries :

- pandas**: for manipulation and analysis of data.(using Dataframes)
- numpy**: to operate on large arrays and matrices with multiple dimensions.
- matplotlib**:to visualize data and get inferences using graphs and other elements.
- seaborn**: for drawing attractive and informative statistical graphics.
- SK-Learn**: for ML Classification Algorithms.
- Sys and Warnings**: to ignore warnings.

(ii) Loading the Dataset

```
In [13]: # Data Importing
df = pd.read_csv("adult.csv")
print(df.shape)
df.head()
```

(32561, 15)

```
Out[13]:
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain	capital.loss	hours.per.week	native.
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female	0	4356	40	United States
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	0	4356	18	United States
2	66	?	186061	Some-college	10	Widowed	?	Unmarried	Black	Female	0	4356	40	United States
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female	0	3900	40	United States
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	0	3900	40	United States

```
In [15]: df.describe()
```

```
Out[15]:
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437456
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347429
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

Dataset Contains Attributes:

Age

Workclass

Final Weight

Education

Education Number of Years

Marital-status.

Occupation

Relationship

Race

Sex

Capital-gain

Capital-loss
Hours-per-week
Native-country

2. Data analysis and exploration

I. Converting missing values to null:

```
# to replace ? with null
change_cols = ['workclass', 'occupation', 'native.country']
for col in change_cols:
    df.loc[df[col] == '?', col] = 'null'

# cross check
for col in change_cols:
    print(f"? in {col}: {df[(df[col] == '?')].any().sum()}")

? in workclass: 0
? in occupation: 0
? in native.country: 0
```

II. Merging and replacing attributes in the list

- **Education:** Merged the High School grad, 1st to 12th into Schooling, Replaced the bachelors to undergraduates and masters to postgraduates.

```
# merging and replacing elements in the list
```

```
school = ['HS-grad', '12th', '11th', '10th', '9th', '1st-4th', '5th-6th', '7th-8th', 'Preschool']  
df['education'].replace(to_replace = school, value = 'Schooling', inplace = True)  
df['education'].replace(to_replace = ['Bachelors'], value = "Undergraduates", inplace = True)  
df['education'].replace(to_replace = ['Masters'], value = "Post-Graduates", inplace = True)  
df['education'].value_counts()
```

```
Schooling      14754  
Some-college   7291  
Undergraduates 5355  
Post-Graduates 1723  
Assoc-voc      1382  
Assoc-acdm     1067  
Prof-school     576  
Doctorate       413  
Name: education, dtype: int64
```

- **Marital Status:** Merged the married-spouse-absent, married-civ-spouse, married-AF-spouse into married, separated and divorced into separated and replaced never-married into single

```
married= ['Married-spouse-absent', 'Married-civ-spouse', 'Married-AF-spouse']  
separated = ['Separated', 'Divorced']  
  
df['marital.status'].replace(to_replace = married, value = 'Married', inplace = True)  
df['marital.status'].replace(to_replace = separated, value = 'Separated', inplace = True)  
df['marital.status'].replace(to_replace = ['Never-married'], value = "Single", inplace = True)  
  
df['marital.status'].value_counts()
```

```
Married      15417  
Single       10683  
Separated     5468  
Widowed        993  
Name: marital.status, dtype: int64
```

- **Work Class:** Merged self-emp-not-inc and self-emp-inc into Self-Employed, merged Local-gov, State-gov and

federal-gov into govt-employees and replaced never worked to unemployed.

```
self_employed = ['Self-emp-not-inc', 'Self-emp-inc']
govt_employees = ['Local-gov', 'State-gov', 'Federal-gov']

df['workclass'].replace(to_replace = self_employed, value = 'Self-Employed', inplace = True)
df['workclass'].replace(to_replace = govt_employees, value = 'Govt-Employees', inplace = True)
df['workclass'].replace(to_replace = ['Never-worked'], value = 'Unemployed', inplace = True)


df['workclass'].value_counts()
```

Private	22696
Govt-Employees	4351
Self-Employed	3657
null	1836
Without-pay	14
Unemployed	7

Name: workclass, dtype: int64


III. Visualization

- Checking the total number of people having income greater than 50K and less than or equal to 50K:



inference: This means almost 75% of people having income less than or equal to 50K and almost 25% of people having income greater than 50K.

- Distribution of population on the basis of numerical features:



inference: 1. Maximum individuals are in age between 15-45
2. Most people work 30-40 hours per week
3. Most individuals studied till 9th standard

- Exploring the data based on Capital Gain and Capital Loss:

***** sex *****

Male 18551

Female 9779

Name: sex, dtype: int64

***** native.country *****

United-States 25320

Mexico 612

null 493

Philippines 174

Germany 117

Puerto-Rico 103

Canada 103

El-Salvador 95

Cuba 85

India 79

Jamaica 78

England 78

South 68

Dominican-Republic 67

Italy 65

China 64

Guatemala 60

Vietnam 57

Columbia 55

Poland 53

Japan 51

Taiwan 44

Haiti 42

Portugal 35

Iran 35

Nicaragua 30

Peru 29

France 26

Ecuador 25

Ireland 21

***** occupation *****

Craft-repair	3593
Adm-clerical	3408
Prof-specialty	3290
Exec-managerial	3219
Sales	3138
Other-service	3122
Machine-op-inspct	1806
null	1662
Transport-moving	1416
Handlers-cleaners	1274
Farming-fishing	890
Tech-support	795
Protective-serv	570
Priv-house-serv	139
Armed-Forces	8

Name: occupation, dtype: int64

***** relationship *****

Husband	10739
Not-in-family	7427
Own-child	4810
Unmarried	3172
Wife	1272
Other-relative	910

Name: relationship, dtype: int64

***** race *****

White	24061
Black	2839
Asian-Pac-Islander	902
Amer-Indian-Eskimo	280
Other	248

Name: race, dtype: int64

```
***** workclass *****
```

```
Private      19982
```

```
Govt-Employees  3714
```

```
Self-Employed  2960
```

```
null         1655
```

```
Without-pay   12
```

```
Unemployed     7
```

```
Name: workclass, dtype: int64
```

```
***** education *****
```

```
Schooling      13342
```

```
Some-college   6533
```

```
Undergraduates 4384
```

```
Post-Graduates 1300
```

```
Assoc-voc      1194
```

```
Assoc-acdm     930
```

```
Prof-school    363
```

```
Doctorate      284
```

```
Name: education, dtype: int64
```

```
***** marital.status *****
```

```
Married        12603
```

```
Single          9914
```

```
Separated       4934
```

```
Widowed         879
```

```
Name: marital.status, dtype: int64
```

inference:

a) 1519 people having capital loss above the median value which is almost 4.67%

b) 2712 people having capital gain above the median value which is almost 8.33%

c) Almost 92% of people having capital gain equals to zero.

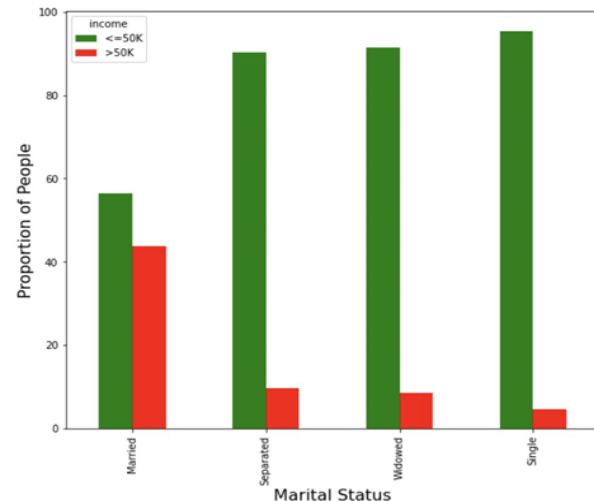

- The maximum capital gain is 99999 by 159 observations






```
Number of observations having capital gain of 99999:(159, 15)
Income counts: >50K    159
Name: income, dtype: int64
```

- The Maximum capital loss is 4356 by 3 observations


```
Number of observations having capital loss of 4356:(3, 15)
Income counts: <=50K    3
Name: income, dtype: int64
```

- comparing between observations having income greater than 50k and less than or equal to 50k on the basis of the categorical attributes:





- Correlation between the attributes



3. Data preprocessing

1.removing the numerical data present in the given data

```
df.drop(['age', 'fnlwgt', 'education.num', 'capital.gain', 'capital.loss', 'hours.per.week'], axis = 1, inplace = True)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 9 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   workclass            32561 non-null  object
1   education            32561 non-null  object
2   marital.status       32561 non-null  object
3   occupation           32561 non-null  object
4   relationship         32561 non-null  object
5   race                32561 non-null  object
6   sex                 32561 non-null  object
7   native.country       32561 non-null  object
8   income              32561 non-null  object
dtypes: object(9)
```

2.changing income into 0 and 1,

0 representing income less than or equal to 50k and

1 representing income greater than 50k

3.Converting Categorical data to numbers to handle them more effectively

	workclass	education	marital.status	occupation	relationship	race	sex	native.country	income
0	4	1	3	13	1	3	1	United-States	0
1	1	1	3	7	1	3	1	United-States	0
2	4	5	3	13	4	0	1	United-States	0
3	1	1	1	5	4	3	1	United-States	0
4	1	5	1	4	0	3	1	United-States	0
...
32556	1	5	2	2	1	3	0	United-States	0
32557	1	2	0	6	2	3	1	United-States	0
32558	1	1	0	5	3	3	0	United-States	1
32559	1	1	3	14	4	3	1	United-States	0
32560	1	1	2	14	0	3	0	United-States	0

4.removing native country attribute as not needed apparently

	workclass	education	marital.status	occupation	relationship	race	sex	income
0	4	1	3	13	1	3	1	0
1	1	1	3	7	1	3	1	0
2	4	5	3	13	4	0	1	0
3	1	1	1	5	4	3	1	0
4	1	5	1	4	0	3	1	0
...
32556	1	5	2	2	1	3	0	0
32557	1	2	0	6	2	3	1	0
32558	1	1	0	5	3	3	0	1
32559	1	1	3	14	4	3	1	0
32560	1	1	2	14	0	3	0	0

5. Now, splitting the data into input features and output label i.e X contains categorical data(predictors) and Y contains income (0/1)

6.Dividing X and Y into training and test sets

```
# Splitting the data in the ratio 3:1 where 3 is for training data and 1 is for testing data
X_train, X_test, Y_train, Y_test = train_test_split(dfx, dfy, test_size = 0.25, random_state = 42)
print(X_train.shape)
print(X_test.shape)
print(Y_train.shape)
print(Y_test.shape)
```

```
(24420, 7)
(8141, 7)
(24420, 1)
(8141, 1)
```

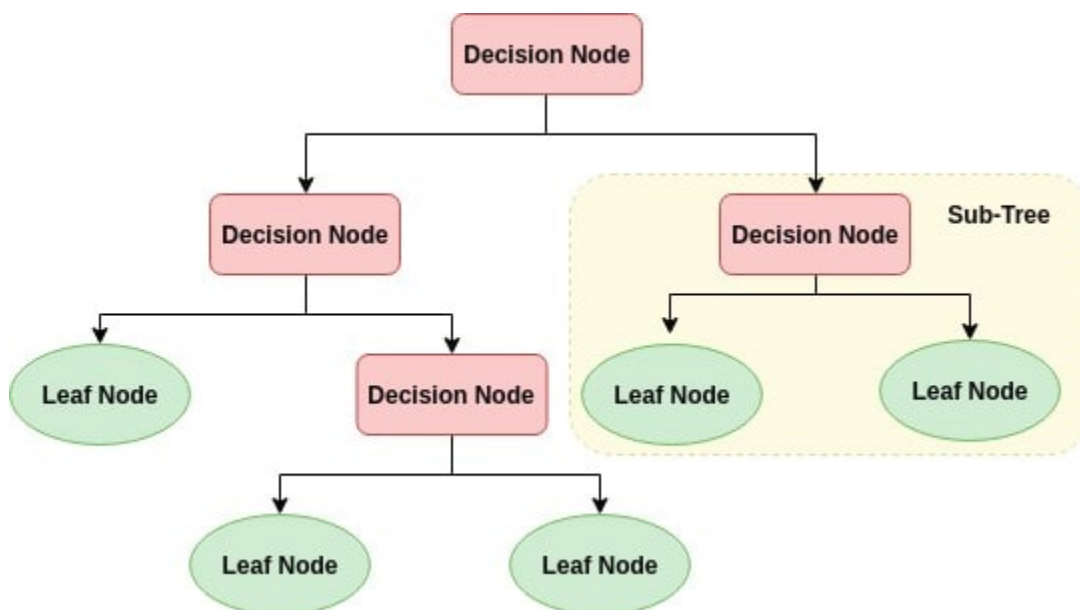
4. Applying ML Classification Algorithms

We would be applying 3 ML algorithms on the given dataset and look for the best one by comparing the accuracy obtained from each of them.

1. Decision Tree Classifier:

In a decision tree, the algorithm begins at the root node and works its way up to predict the class of a given dataset. This algorithm checks the values of the root attribute with the values of the record (actual dataset) attribute and then follows the branch and jumps to the next node based on the comparison.

The algorithm compares the attribute value with the other sub-nodes and moves on to the next node. It repeats the process until it reaches the tree's leaf node.

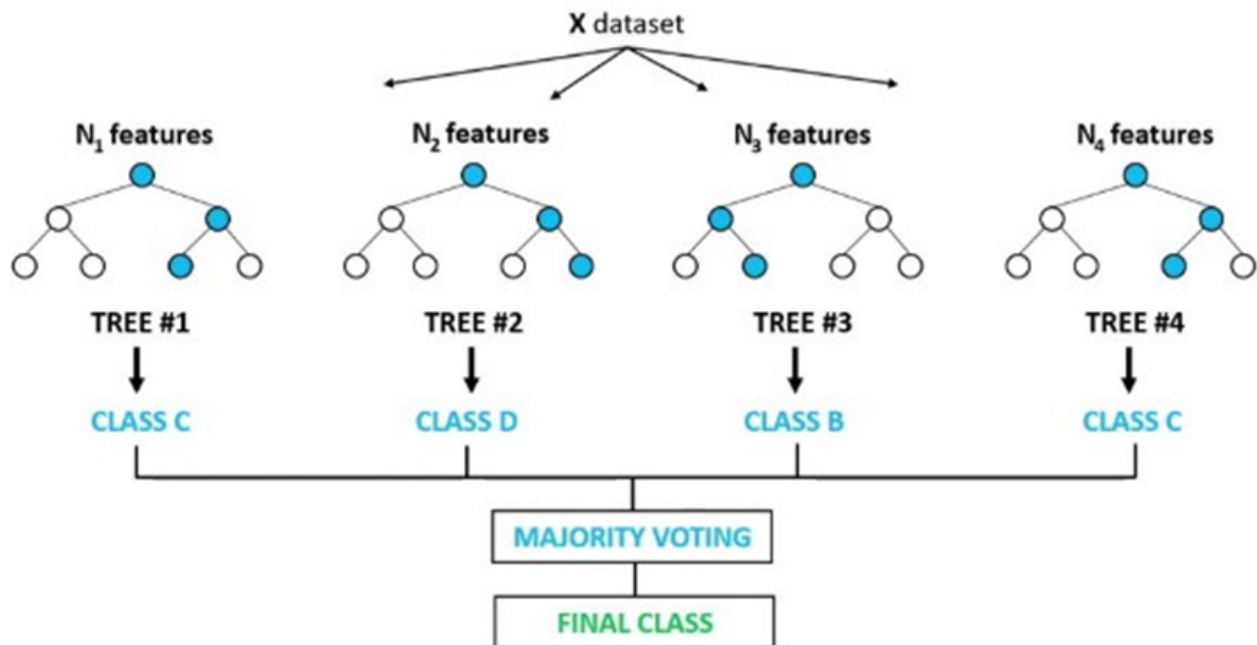


2. Random Forest Classifier:

This process is based on learning together. Many decision trees are not a single tree but work as a whole. It increases the randomness of the integration by randomly generating a forest of decision trees. Each tree decides to produce the output, and the final ordering of the input data is determined by majority vote.


They are more accurate than single decision trees but take longer to train.

Random Forest Classifier



3. Support Vector Machine

Support vector machines use supervised learning algorithms. It is used to find hyperplanes separating groups of data. Now there may be many planes separating the data, but SVM tries to find the best line for this separation. When there is a clear dividing line between groups and the data sets are not large enough.



5.Implementing these algorithms and displaying their confusion matrix

1.Decision Tree Classifier

Accuracy: 82%

Right Prediction: $5661 + 1015 = 6676$

Wrong Prediction: $929 + 536 = 1492$


Decision Tree Classifier

```
operation2 = DecisionTreeClassifier()  
operation2.fit(X_train, Y_train)  
Y_pred = operation2.predict(X_test)
```

```
score2 = accuracy_score(Y_test, Y_pred)  
print('Prediction Accuracy = ' + str(score2*100))
```

Prediction Accuracy = 82.00466773123695

```
class_names = [0,1]  
fig, ax = plt.subplots(figsize=(8,4))  
plot_confusion_matrix(operation2, X_test, Y_test, cmap=plt.cm.Reds, labels=class_names, ax=ax, values_format = '.0f')  
plt.title('Confusion Matrix')  
plt.grid(False)  
plt.show()
```



2.Random Forest Classifier

Accuracy: 82.36%

Right Prediction: $5662 + 535 = 6705$

Wrong Prediction: $901 + 535 = 1436$


Random Forest ¶

```
operation4 = RandomForestClassifier()  
operation4.fit(X_train, Y_train)  
Y_pred = operation4.predict(X_test)
```

```
score4 = accuracy_score(Y_test, Y_pred)  
print('Prediction Accuracy = ' + str(score4*100))
```

Prediction Accuracy = 82.36088932563567

```
class_names = [0,1]  
fig, ax = plt.subplots(figsize=(8,4))  
plot_confusion_matrix(operation4, X_test, Y_test, cmap=plt.cm.Red, labels=class_names, ax=ax, values_format = '.0f')  
plt.title('Confusion Matrix')  
plt.grid(False)  
plt.show()
```



3.Support Vector Machine

Accuracy: 80.75%

Right Prediction: $5719 + 855 = 6574$

Wrong Prediction: $1089 + 478 = 1567$


Support Vector Machine

```
operation5 = SVC()  
operation5.fit(X_train, Y_train)  
Y_pred = operation5.predict(X_test)
```

```
score5 = accuracy_score(Y_test, Y_pred)  
print('Prediction Accuracy = ' + str(score5*100))
```

Prediction Accuracy = 80.75175039921385

```
class_names = [0,1]  
fig, ax = plt.subplots(figsize=(8,4))  
plot_confusion_matrix(operation5, X_test, Y_test, cmap=plt.cm.Red, labels=class_names, ax=ax, values_format = '.0f')  
plt.title('Confusion Matrix')  
plt.grid(False)  
plt.show()
```



6. Conclusion

Algorithm	Accuracy
Decision Tree	82%
Random Forest	82.36%
SVM	80.75%

Thus after performing data collection, data exploration , data preprocessing and finally training 3 different models for the given problem we conclude that Random Forest is giving the best accuracy and is the most suited algorithm here.

7.References

- Class notes and discussions
- Documentation of python libraries i.e
Matplotlib,seaborn,scikit-learn,pandas and numpy
- <https://archive.ics.uci.edu/ml/datasets/adult>